

Teja Gollapudi

teja@tejagollapudi.com * 3126841770 * Palo Alto, CA

[GitHub](#) * [Medium](#) * [Website](#) * [LinkedIn](#)

Work Experience

Meta

Research Engineer

Jan 2024 – Present

Burlingame, CA

Post-Training for Foundation Models

- Drive agentic improvements to Muse for wearables using reinforcement learning with verifiable rewards (RLVR), with a focus on long-horizon, multi-step task execution
- Strengthen agentic behavior in foundation models through reinforcement learning, advancing tool calling across multi-turn, multi-tool, and multimodal settings
- Built synthetic data generation pipelines across multiple projects, covering tasks from single-turn tool calling and summarization to long-horizon agentic workflows
- Improve LLM reasoning and factuality; published findings at internal and external venues
- Led end-to-end development of multiple wearables AI models, from design and training through evaluation and production deployment

Wearables Product Launches

- Helped launch Muse on wearables, contributing to RL training, synthetic data generation, and evaluation
- Led real-time AI sports features on wearables for Super Bowl LIX, the Olympics, and the FIFA Club World Cup
- Enabled AI-native product launches, including Meta Ray-Ban Display and Oakley Meta sports glasses
- Led development of AI fitness features that surface workout history on wearables

VMware Inc.

Machine Learning Engineer

2021 - 2024

Palo Alto, CA

Large Language Models (LLMs)

- Co-led development of commercially viable instruction-following open-source LLMs, publishing models on [Hugging Face](#) and [GitHub](#)
- Fine-tuned code LLM for an internal code completion co-pilot service
- Designed experiments, data-curation strategies, and fine-tuning pipelines for both low-resource hardware and multi-node cloud GPU training
- Provided technical guidance to teams across the organization on adopting LLMs

LLM API

- Built a self-contained internal LLM API service and Python client, handling **2,000+** daily requests across **40+** teams
- Scaled production systems with vLLM and Ray for enhanced throughput

NLP Models and Optimization

- Fine-tuned [RoBERTa](#) and distilled into [TinyRoBERTa](#), achieving **600%** throughput increase
- Trained multilingual ColBERT model achieving **7x** higher Recall@5 over previous search solution
- Pre-trained [domain-specific BERT](#) and distilled into [MiniLMv2](#), open-sourcing both [libraries](#)

Leadership and Community

- Founded [VMware's Hugging Face page](#), driving monthly downloads into the **five-digit** range
- Published [technical blogs](#) and papers at internal and external venues
- Mentored engineers transitioning into ML roles; conducted interviews for ML positions

Machine Learning Engineering Intern

VMware, Inc.

Jun 2020 - Jan 2021

Palo Alto, CA

- Created an ML spell checker using Seq2Seq transformer architecture

Machine Learning Intern

Inkers.ai

May 2018 - Jul 2018

Bengaluru, India

- Designed a tree-structured Convolutional Neural Network for object detection using Multi-task learning
- Constructed a large image dataset to train multiple CNN architectures

Technical skills

Domains	Machine Learning, Natural Language Processing, Generative Models, Computer Vision, Reinforcement Learning
ML Libraries	PyTorch, TorchScript, Tensorflow, Keras, Hugging Face libraries, Numpy, Deepspeed, ONNX Runtime
Languages	Python (Preferred), Java, C++ (Familiar), Javascript (Familiar)
Other	Docker, Kubernetes, FastAPI, SQL, MongoDB, Cloud Platforms (GCP, AWS S3), Ray

Education

Master of Science in Computer Science

Aug 2019 - May 2021

University of Illinois at Chicago (UIC)

4.0 GPA

B.Tech in Information Technology

Jul 2015 - May 2019

Vellore Institute of Technology (VIT), India

Publications

-
- [TruthRL: Incentivizing Truthful LLMs via Reinforcement Learning](#), ICML 2026
 - [PrismRAG: Boosting RAG Factuality with Distractor Resilience and Strategized Reasoning](#), M. Kachuee, **T. Gollapudi**, EMNLP 2025
 - [The Unreasonable Effectiveness of Eccentric Automatic Prompts](#), R. Battle, **T. Gollapudi**, arXiv 2024
 - [Fast Multi-scale Face Detection CNN](#), **N. S. Gollapudi**, ViTECoN 2019

Miscellaneous

-
- **Press:** Work featured in [IEEE Spectrum](#), [The Register](#), [Business Insider](#), and more